

untitled1

June 26, 2024

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[5]: data=pd.read_csv('movie_metadata.csv')
data.head()
```

```
[5]:
```

	color	director_name	num_critic_for_reviews	duration	\
0	Color	James Cameron	723.0	178.0	
1	Color	Gore Verbinski	302.0	169.0	
2	Color	Sam Mendes	602.0	148.0	
3	Color	Christopher Nolan	813.0	164.0	
4	NaN	Doug Walker	NaN	NaN	

	director_facebook_likes	actor_3_facebook_likes	actor_2_name	\
0	0.0	855.0	Joel David Moore	
1	563.0	1000.0	Orlando Bloom	
2	0.0	161.0	Rory Kinnear	
3	22000.0	23000.0	Christian Bale	
4	131.0	NaN	Rob Walker	

	actor_1_facebook_likes	gross	genres	...	\
0	1000.0	760505847.0	Action Adventure Fantasy Sci-Fi	...	
1	40000.0	309404152.0	Action Adventure Fantasy	...	
2	11000.0	200074175.0	Action Adventure Thriller	...	
3	27000.0	448130642.0	Action Thriller	...	
4	131.0	NaN	Documentary	...	

	num_user_for_reviews	language	country	content_rating	budget	\
0	3054.0	English	USA	PG-13	237000000.0	
1	1238.0	English	USA	PG-13	300000000.0	
2	994.0	English	UK	PG-13	245000000.0	
3	2701.0	English	USA	PG-13	250000000.0	
4	NaN	NaN	NaN	NaN	NaN	

	title_year	actor_2_facebook_likes	imdb_score	aspect_ratio	\
--	------------	------------------------	------------	--------------	---

0	2009.0	936.0	7.9	1.78
1	2007.0	5000.0	7.1	2.35
2	2015.0	393.0	6.8	2.35
3	2012.0	23000.0	8.5	2.35
4	NaN	12.0	7.1	NaN

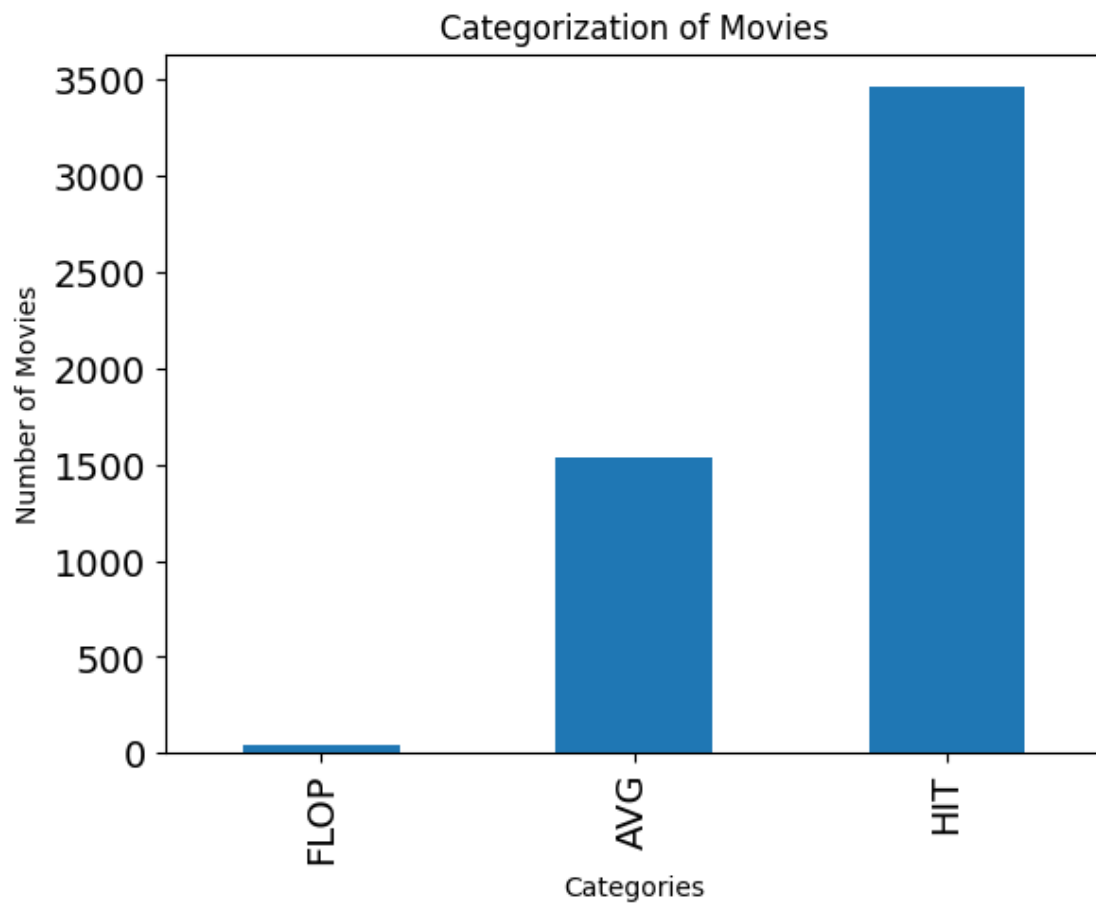
	movie_facebook_likes
0	33000
1	0
2	85000
3	164000
4	0

[5 rows x 28 columns]

```
[6]: bins = [ 1, 3, 6, 10]
      labels = ['FLOP', 'AVG', 'HIT']
      data['imdb_binned'] = pd.cut(data['imdb_score'], bins=bins, labels=labels)
```

```
[7]: data.groupby(['imdb_binned']).size().plot(kind="bar",fontsize=14)
      plt.xlabel('Categories')
      plt.ylabel('Number of Movies')
      plt.title('Categorization of Movies')
```

```
[7]: Text(0.5, 1.0, 'Categorization of Movies')
```



```
[8]: data.head(5)
```

```
[8]:   color      director_name  num_critic_for_reviews  duration  \
0  Color      James Cameron                723.0    178.0
1  Color      Gore Verbinski                302.0    169.0
2  Color          Sam Mendes                602.0    148.0
3  Color  Christopher Nolan                813.0    164.0
4   NaN          Doug Walker                 NaN      NaN

   director_facebook_likes  actor_3_facebook_likes  actor_2_name  \
0                   0.0                855.0  Joel David Moore
1                 563.0                1000.0   Orlando Bloom
2                   0.0                 161.0    Rory Kinnear
3                22000.0               23000.0  Christian Bale
4                 131.0                  NaN    Rob Walker

   actor_1_facebook_likes  gross  genres  ...  \
0                 1000.0  760505847.0  Action|Adventure|Fantasy|Sci-Fi  ...
```

1	40000.0	309404152.0	Action Adventure Fantasy	...
2	11000.0	200074175.0	Action Adventure Thriller	...
3	27000.0	448130642.0	Action Thriller	...
4	131.0	NaN	Documentary	...

	language	country	content_rating	budget	title_year	\
0	English	USA	PG-13	2370000000.0	2009.0	
1	English	USA	PG-13	3000000000.0	2007.0	
2	English	UK	PG-13	2450000000.0	2015.0	
3	English	USA	PG-13	2500000000.0	2012.0	
4	NaN	NaN	NaN	NaN	NaN	

	actor_2_facebook_likes	imdb_score	aspect_ratio	movie_facebook_likes	\
0	936.0	7.9	1.78	33000	
1	5000.0	7.1	2.35	0	
2	393.0	6.8	2.35	85000	
3	23000.0	8.5	2.35	164000	
4	12.0	7.1	NaN	0	

	imdb_binned
0	HIT
1	HIT
2	HIT
3	HIT
4	HIT

[5 rows x 29 columns]

```
[9]: data.shape
```

```
[9]: (5043, 29)
```

```
[10]: data.isnull().sum()
```

```
[10]: color                19
      director_name       104
      num_critic_for_reviews  50
      duration             15
      director_facebook_likes 104
      actor_3_facebook_likes  23
      actor_2_name           13
      actor_1_facebook_likes   7
      gross                 884
      genres                 0
      actor_1_name            7
      movie_title             0
      num_voted_users          0
```

```

cast_total_facebook_likes    0
actor_3_name                 23
facenumber_in_poster        13
plot_keywords                153
movie_imdb_link              0
num_user_for_reviews         21
language                    14
country                     5
content_rating               303
budget                      492
title_year                  108
actor_2_facebook_likes       13
imdb_score                   0
aspect_ratio                 329
movie_facebook_likes         0
imdb_binned                  0
dtype: int64

```

```
[11]: data.dropna(inplace=True)
```

```
[12]: data.shape
```

```
[12]: (3755, 29)
```

```
[13]: data.columns
```

```
[13]: Index(['color', 'director_name', 'num_critic_for_reviews', 'duration',
'director_facebook_likes', 'actor_3_facebook_likes', 'actor_2_name',
'actor_1_facebook_likes', 'gross', 'genres', 'actor_1_name',
'movie_title', 'num_voted_users', 'cast_total_facebook_likes',
'actor_3_name', 'facenumber_in_poster', 'plot_keywords',
'movie_imdb_link', 'num_user_for_reviews', 'language', 'country',
'content_rating', 'budget', 'title_year', 'actor_2_facebook_likes',
'imdb_score', 'aspect_ratio', 'movie_facebook_likes', 'imdb_binned'],
dtype='object')
```

```
[14]: data.shape
```

```
[14]: (3755, 29)
```

```
[15]: data.describe(include='object')
```

```
[15]:
```

	color	director_name	actor_2_name	genres \
count	3755	3755	3755	3755
unique	2	1658	2187	745
top	Color	Steven Spielberg	Morgan Freeman	Comedy Drama Romance
freq	3631	25	20	147

	actor_1_name	movie_title	actor_3_name	\
count	3755	3755	3755	
unique	1427	3654	2586	
top	Robert De Niro	Victor Frankenstein	Steve Coogan	
freq	42	3	8	

	plot_keywords	\
count	3755	
unique	3655	
top	halloween masked killer michael myers slasher ...	
freq	3	

	movie_imdb_link	language	country	\
count	3755	3755	3755	
unique	3655	33	45	
top	http://www.imdb.com/title/tt0077651/?ref_=fn_t...	English	USA	
freq	3	3598	2986	

	content_rating
count	3755
unique	12
top	R
freq	1700

```
[16]: data.drop(columns=['movie_title', 'movie_imdb_link'], inplace=True)
```

```
[18]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
cat_list=['color', 'director_name', 'actor_2_name',
          'genres', 'actor_1_name',
          'actor_3_name',
          'plot_keywords',
          'language', 'country', 'content_rating',
          'title_year', 'aspect_ratio']
data[cat_list]=data[cat_list].apply(lambda x:le.fit_transform(x))
```

```
[19]: data.head()
```

	color	director_name	num_critic_for_reviews	duration	\
0	1	620	723.0	178.0	
1	1	538	302.0	169.0	
2	1	1394	602.0	148.0	
3	1	251	813.0	164.0	
5	1	62	462.0	132.0	

	director_facebook_likes	actor_3_facebook_likes	actor_2_name	\
--	-------------------------	------------------------	--------------	---

0	0.0	855.0	1001
1	563.0	1000.0	1591
2	0.0	161.0	1794
3	22000.0	23000.0	380
5	475.0	530.0	1836

	actor_1_facebook_likes	gross	genres	...	language	country	\
0	1000.0	760505847.0	91	...	9	43	
1	40000.0	309404152.0	85	...	9	43	
2	11000.0	200074175.0	107	...	9	42	
3	27000.0	448130642.0	243	...	9	43	
5	640.0	73058679.0	105	...	9	43	

	content_rating	budget	title_year	actor_2_facebook_likes	\
0	7	237000000.0	66	936.0	
1	7	300000000.0	64	5000.0	
2	7	245000000.0	72	393.0	
3	7	250000000.0	69	23000.0	
5	7	263700000.0	69	632.0	

	imdb_score	aspect_ratio	movie_facebook_likes	imdb_binned
0	7.9	7	33000	HIT
1	7.1	12	0	HIT
2	6.8	12	85000	HIT
3	8.5	12	164000	HIT
5	6.6	12	24000	HIT

[5 rows x 27 columns]

```
[24]: numeric_data = data.select_dtypes(include=np.number)

# Compute the correlation matrix
corr = numeric_data.corr()


# Create a mask to hide the upper triangle of the plot
mask = np.zeros_like(corr, dtype=bool)
mask[np.triu_indices_from(mask)] = True

# Setting up the matplotlib figure
plt.figure(figsize=(20, 15))

# Plotting the heatmap using seaborn
sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns,
            cmap='RdYlGn', annot=True, mask=mask)

# Adding title to the plot
plt.title('Correlation Heatmap')
```

```
# Displaying the plot
plt.show()
```



```
[25]: data.
      drop(columns=['cast_total_facebook_likes', 'num_critic_for_reviews'], inplace=True)
```

```
[26]: data.drop(columns=['imdb_score'], inplace=True)
```

```
[27]: data.shape
```

```
[27]: (3755, 24)
```

```
[30]: X = data.iloc[:, 0:23].values
      #Dependent/Target Variable
      y = data.iloc[:, 23].values
      y
```



```
[30]: ['HIT', 'HIT', 'HIT', 'HIT', 'HIT', ..., 'HIT', 'HIT', 'HIT', 'HIT', 'HIT']
Length: 3755
Categories (3, object): ['FLOP' < 'AVG' < 'HIT']
```

```
[32]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
    random_state = 0, stratify = y)
print(X_train.shape)
print(y_train.shape)
```

```
(2628, 23)
```

```
(2628,)
```

```
[33]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
[34]: from sklearn.feature_selection import RFECV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import log_loss
clf_rf=RandomForestClassifier(random_state=0)
rfecv=RFECV(estimator=clf_rf, step=1,cv=5,scoring='neg_log_loss')
rfecv=rfecv.fit(X_train,y_train)
```

```
[35]: X_train = pd.DataFrame(X_train)
X_test = pd.DataFrame(X_test)
print('Optimal number of features :', rfecv.n_features_)
print('Best features :', X_train.columns[rfecv.support_])
```

```
Optimal number of features : 22
```

```
Best features : Index([1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16,
17, 18, 19, 20,
21, 22],
dtype='int64')
```

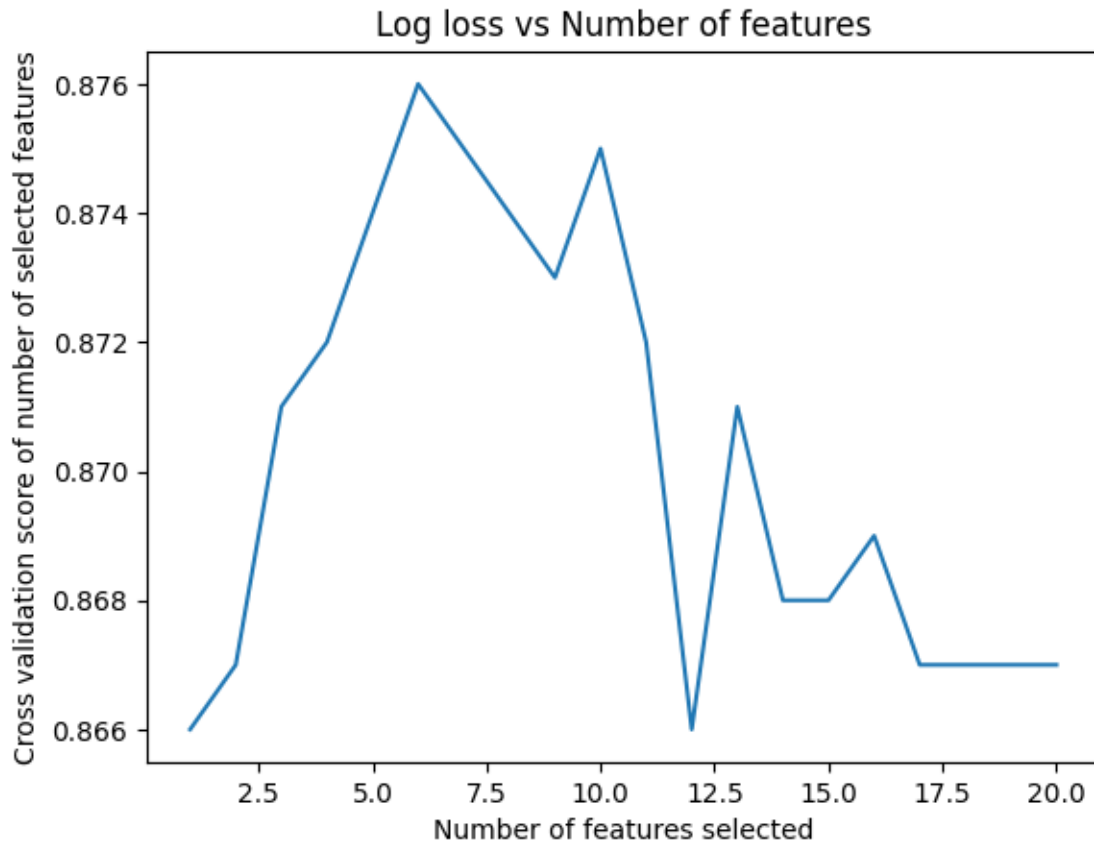
```
[36]: clf_rf = clf_rf.fit(X_train,y_train)
importances = clf_rf.feature_importances_

std = np.std([tree.feature_importances_ for tree in clf_rf.estimators_],
axis=0)
indices = np.argsort(importances)[::-1]
```

```
[63]: import matplotlib.pyplot as plt

plt.figure()
plt.xlabel("Number of features selected")
```

```
plt.ylabel("Cross validation score of number of selected features")
plt.title("Log loss vs Number of features")
plt.plot(range(1, len(rfecv.cv_results_['mean_test_score']) + 1), rfecv.
         ↪cv_results_['mean_test_score'])
plt.show()
```



```
[70]: # Ensure that we are working with the data1 containing the correct features
data1 = data.copy()

# Drop the 'imdb_binned' column if it exists
if 'imdb_binned' in data1.columns:
    data1.drop(columns=['imdb_binned'], inplace=True)

# Selecting the Important Features
X_opt_train = rfecv.transform(X_train)
X_opt_test = rfecv.transform(X_test)

# Scaling the selected features
sc = StandardScaler()
X_opt_train = sc.fit_transform(X_opt_train)
```

```

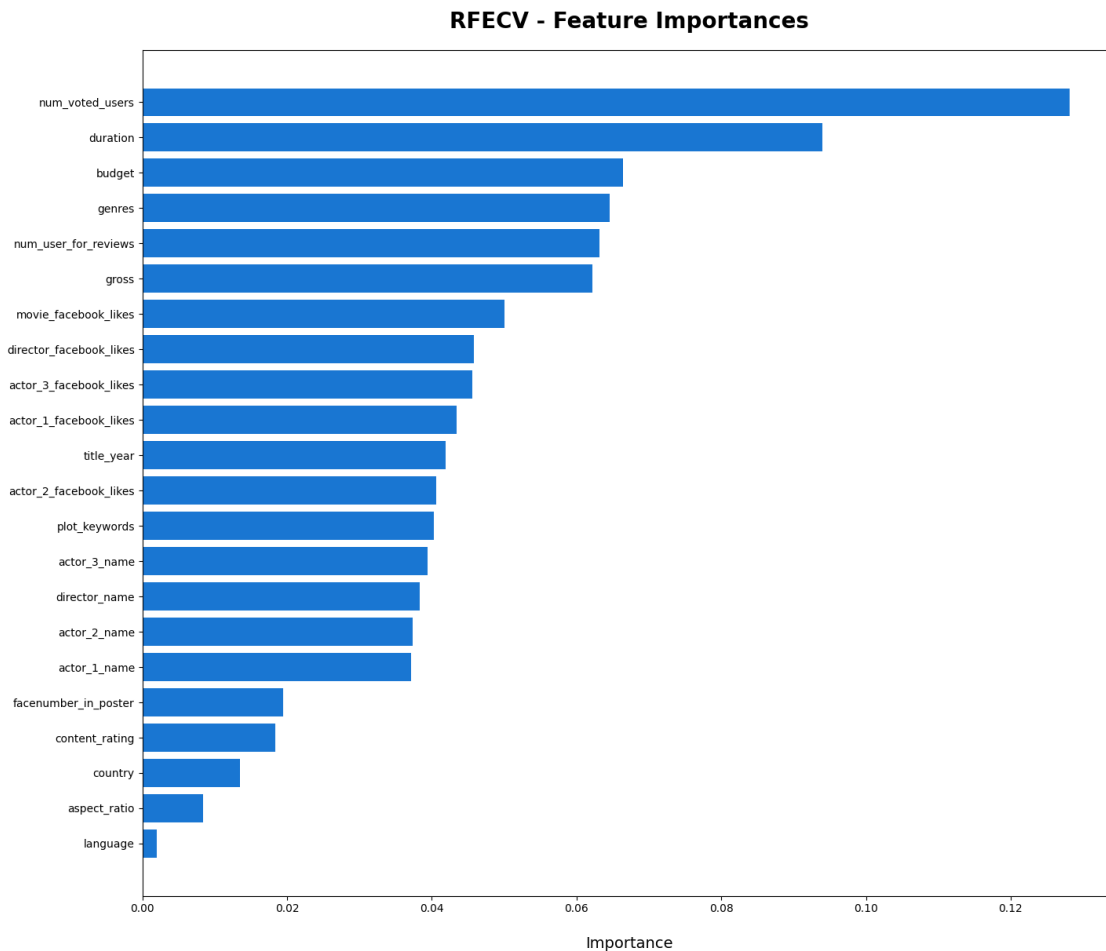
X_opt_test = sc.transform(X_opt_test)

# Creating a new dataframe with column names and feature importance
dset = pd.DataFrame()
dset['attr'] = data1.columns[rfecv.support_]
dset['importance'] = rfecv.estimator_.feature_importances_

# Sorting with importance column
dset = dset.sort_values(by='importance', ascending=True)

# Barplot indicating Feature Importance
plt.figure(figsize=(16, 14))
plt.barh(y=dset['attr'], width=dset['importance'], color='#1976D2')
plt.title('RFECV - Feature Importances', fontsize=20, fontweight='bold', pad=20)
plt.xlabel('Importance', fontsize=14, labelpad=20)
plt.show()

```



```
[90]: from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(n_estimators = 100, criterion = 'entropy',
    ↪ random_state = 0)
classifier.fit(X_opt_train,y_train)
```

```
[90]: RandomForestClassifier(criterion='entropy', random_state=0)
```

```
[92]: y_pred = classifier.predict(X_opt_test)
```

```
[93]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test,y_pred)
cm
```

```
[93]: array([[191,  0, 144],
           [  4,  0,  5],
           [ 68,  0, 715]])
```

```
[94]: from sklearn.metrics import classification_report
cr = classification_report(y_test,y_pred)
print(cr)
```

	precision	recall	f1-score	support
AVG	0.73	0.57	0.64	335
FLOP	0.00	0.00	0.00	9
HIT	0.83	0.91	0.87	783
accuracy			0.80	1127
macro avg	0.52	0.49	0.50	1127
weighted avg	0.79	0.80	0.79	1127

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```